

27. EDA - Income Data

October 31, 2022

1 Exploratory Data Analysis on Income Data

Data-set `income_evaluation.csv` was extracted from the 1994 Census bureau database of USA.

Columns/Features:

1. age: continuous
2. workclass: categorical
3. fnlwgt: continuous
4. education: categorical
5. education-num: continuous
6. marital-status: categorical
7. occupation: categorical
8. relationship: categorical
9. race: categorical
10. sex: categorical
11. capital-gain: continuous
12. capital-loss: continuous
13. hours-per-week: continuous
14. native-country: categorical
15. income: target

2 Import required packages

```
[1]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import seaborn as sns
```

3 Load dataset

```
[2]: df = pd.read_csv('../data/income_evaluation.csv')
```

4 Exploratory Data Analysis

Explore the data to gain insights about the data.

4.1 View dimensions of dataset

```
[3]: df.shape
```

```
[3]: (32561, 15)
```

We can see that there are 32561 instances and 15 attributes in the data set.

4.2 Preview the dataset

```
[4]: df.head()
```

```
[4]:
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

4.3 Rename column names

We can see that the dataset does not have proper column names. The column names contain underscore. We should give proper names to the columns. I will do it as follows:-

```
[5]: df.columns
```

```
[5]: Index(['age', ' workclass', ' fnlwgt', ' education', ' education-num',
        ' marital-status', ' occupation', ' relationship', ' race', ' sex',
        ' capital-gain', ' capital-loss', ' hours-per-week', ' native-country',
        ' income'],
        dtype='object')
```

```
[6]: df.columns = [i.replace('-', '_').strip() for i in df.columns]
```

```
[7]: df.columns
```

```
[7]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
        'marital_status', 'occupation', 'relationship', 'race', 'sex',
        'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
        'income'],
        dtype='object')
```

4.4 View summary of dataset

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education             32561 non-null  object
4   education_num         32561 non-null  int64
5   marital_status        32561 non-null  object
6   occupation            32561 non-null  object
7   relationship          32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital_gain          32561 non-null  int64
11  capital_loss          32561 non-null  int64
12  hours_per_week        32561 non-null  int64
13  native_country        32561 non-null  object
14  income                32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

Findings

- We can see that the dataset contains 9 character variables and 6 numerical variables.
- There are no missing values in the dataset.

4.5 Check the data types of columns

- The above `df.info()` command gives us the number of filled values along with the data types of columns.
- If we simply want to check the data type of a particular column, we can use the following command.

```
[9]: df.dtypes
```

```
[9]: age                int64
workclass             object
fnlwgt               int64
education            object
education_num        int64
marital_status       object
occupation           object
relationship         object
race                object
sex                 object
capital_gain         int64
capital_loss         int64
hours_per_week       int64
native_country       object
income              object
dtype: object
```

4.6 View statistical properties of dataset

```
[10]: df.describe()
```

```
[10]:
```

	age	fnlwgt	education_num	capital_gain	capital_loss	\
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	

	hours_per_week
count	32561.000000
mean	40.437456
std	12.347429
min	1.000000
25%	40.000000
50%	40.000000

```

75%          45.000000
max          99.000000

```

- The above `df.describe()` command presents statistical properties in vertical form.
- If we want to view the statistical properties in horizontal form, we should run the following command.

```
[11]: df.describe().T
```

```

[11]:          count          mean          std          min          25%  \
age          32561.0          38.581647          13.640433          17.0          28.0
fnlwgt        32561.0      189778.366512      105549.977697      12285.0      117827.0
education_num  32561.0          10.080679          2.572720           1.0           9.0
capital_gain   32561.0       1077.648844       7385.292085           0.0           0.0
capital_loss   32561.0          87.303830         402.960219           0.0           0.0
hours_per_week 32561.0          40.437456          12.347429           1.0          40.0

          50%          75%          max
age          37.0          48.0          90.0
fnlwgt       178356.0      237051.0      1484705.0
education_num  10.0          12.0          16.0
capital_gain    0.0           0.0       99999.0
capital_loss    0.0           0.0        4356.0
hours_per_week  40.0          45.0          99.0

```

```
[12]: df.describe(include='all')
```

```

[12]:          age workclass          fnlwgt education education_num  \
count      32561.000000          32561      3.256100e+04          32561      32561.000000
unique           NaN              9              NaN              16              NaN
top            NaN      Private              NaN      HS-grad              NaN
freq           NaN       22696              NaN       10501              NaN
mean          38.581647           NaN      1.897784e+05           NaN       10.080679
std           13.640433           NaN      1.055500e+05           NaN       2.572720
min           17.000000           NaN      1.228500e+04           NaN       1.000000
25%           28.000000           NaN      1.178270e+05           NaN       9.000000
50%           37.000000           NaN      1.783560e+05           NaN      10.000000
75%           48.000000           NaN      2.370510e+05           NaN      12.000000
max           90.000000           NaN      1.484705e+06           NaN      16.000000

          marital_status          occupation relationship          race          sex  \
count              32561              32561          32561      32561      32561
unique                7              15              6              5              2
top      Married-civ-spouse      Prof-specialty      Husband      White      Male
freq              14976              4140          13193      27816      21790
mean                NaN              NaN              NaN              NaN              NaN
std                NaN              NaN              NaN              NaN              NaN

```

min	NaN	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN	NaN

	capital_gain	capital_loss	hours_per_week	native_country	income
count	32561.000000	32561.000000	32561.000000	32561	32561
unique	NaN	NaN	NaN	42	2
top	NaN	NaN	NaN	United-States	<=50K
freq	NaN	NaN	NaN	29170	24720
mean	1077.648844	87.303830	40.437456	NaN	NaN
std	7385.292085	402.960219	12.347429	NaN	NaN
min	0.000000	0.000000	1.000000	NaN	NaN
25%	0.000000	0.000000	40.000000	NaN	NaN
50%	0.000000	0.000000	40.000000	NaN	NaN
75%	0.000000	0.000000	45.000000	NaN	NaN
max	99999.000000	4356.000000	99.000000	NaN	NaN

4.7 Check for missing values

- In Python missing data is represented by two values:
 - **None** : None is a Python singleton object that is often used for missing data in Python code.
 - **NaN** : NaN is an acronym for Not a Number. It is a special floating-point value recognized by all systems that use the standard IEEE floating-point representation.
- There are different methods in place on how to detect missing values.

Pandas `isnull()` and `notnull()` functions

- Pandas offers two functions to test for missing values - **`isnull()`** and **`notnull()`**.
- These are simple functions that return a boolean value indicating whether the passed in argument value is in fact missing data.

Below, I will list some useful commands to deal with missing values.

Useful commands to detect missing values

- **`df.isnull()`**

The above command checks whether each cell in a dataframe contains missing values or not. If the cell contains missing value, it returns True otherwise it returns False.

- **`df.isnull().sum()`**

The above command returns total number of missing values in each column in the dataframe.

- **`df.isnull().sum().sum()`**

It returns total number of missing values in the dataframe.

- `df.isnull().mean()`

It returns percentage of missing values in each column in the dataframe.

- `df.isnull().any()`

It checks which column has null values and which has not. The columns which has null values returns TRUE and FALSE otherwise.

- `df.isnull().any().any()`

It returns a boolean value indicating whether the dataframe has missing values or not. If dataframe contains missing values it returns TRUE and FALSE otherwise.

- `df.isnull().values.any()`

It checks whether a particular column has missing values or not. If the column contains missing values, then it returns TRUE otherwise FALSE.

- `df.isnull().values.sum()`

It returns the total number of missing values in the dataframe.

```
[13]: # check for missing values
df.isnull().sum()
```

```
[13]: age                0
workclass              0
fnlwgt                0
education              0
education_num          0
marital_status         0
occupation             0
relationship           0
race                  0
sex                   0
capital_gain           0
capital_loss           0
hours_per_week         0
native_country         0
income                0
dtype: int64
```

Interpretation

We can see that there are no missing values in the dataset.

4.7.1 Check with assert statement

- We must confirm that our dataset has no missing values.
- We can write an **Assert statement** to verify this.

- We can use an assert statement to programmatically check that no missing, unexpected 0 or negative values are present.
- This gives us confidence that our code is running properly.
- **Assert statement** will return nothing if the value being tested is true and will throw an AssertionError if the value is false.
- Asserts
 - `assert 1 == 1` (return Nothing if the value is True)
 - `assert 1 == 2` (return AssertionError if the value is False)

```
[14]: #assert that there are no missing values in the dataframe

assert pd.notnull(df).all().all()
```

Interpretation

- The above command does not throw any error. Hence, it is confirmed that there are no missing or negative values in the dataset.
- All the values are greater than or equal to zero excluding character values.

4.8 Explore Categorical Variables

```
[15]: categorical = ['workclass', 'education', 'marital_status',
                    'occupation', 'relationship', 'race', 'sex',
                    'native_country']
```

```
[16]: df[categorical].head()
```

```
[16]:
```

	workclass	education	marital_status	occupation \
0	State-gov	Bachelors	Never-married	Adm-clerical
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial
2	Private	HS-grad	Divorced	Handlers-cleaners
3	Private	11th	Married-civ-spouse	Handlers-cleaners
4	Private	Bachelors	Married-civ-spouse	Prof-specialty

	relationship	race	sex	native_country
0	Not-in-family	White	Male	United-States
1	Husband	White	Male	United-States
2	Not-in-family	White	Male	United-States
3	Husband	Black	Male	United-States
4	Wife	Black	Female	Cuba

4.8.1 Frequency distribution of categorical variables

Now, we will check the frequency distribution of categorical variables.


```
[17]: for var in categorical:
      print(df[var].value_counts(), '\n')
```

```
Private          22696
Self-emp-not-inc  2541
Local-gov        2093
?                1836
State-gov        1298
Self-emp-inc     1116
Federal-gov      960
Without-pay      14
Never-worked     7
Name: workclass, dtype: int64
```

```
HS-grad          10501
Some-college     7291
Bachelors        5355
Masters          1723
Assoc-voc        1382
11th             1175
Assoc-acdm       1067
10th             933
7th-8th          646
Prof-school      576
9th              514
12th             433
Doctorate        413
5th-6th          333
1st-4th          168
Preschool        51
Name: education, dtype: int64
```

```
Married-civ-spouse  14976
Never-married      10683
Divorced            4443
Separated           1025
Widowed             993
Married-spouse-absent  418
Married-AF-spouse    23
Name: marital_status, dtype: int64
```

```
Prof-specialty    4140
Craft-repair      4099
Exec-managerial   4066
Adm-clerical      3770
Sales             3650
Other-service     3295
```

Machine-op-inspct	2002
?	1843
Transport-moving	1597
Handlers-cleaners	1370
Farming-fishing	994
Tech-support	928
Protective-serv	649
Priv-house-serv	149
Armed-Forces	9

Name: occupation, dtype: int64

Husband	13193
Not-in-family	8305
Own-child	5068
Unmarried	3446
Wife	1568
Other-relative	981

Name: relationship, dtype: int64

White	27816
Black	3124
Asian-Pac-Islander	1039
Amer-Indian-Eskimo	311
Other	271

Name: race, dtype: int64

Male	21790
Female	10771

Name: sex, dtype: int64

United-States	29170
Mexico	643
?	583
Philippines	198
Germany	137
Canada	121
Puerto-Rico	114
El-Salvador	106
India	100
Cuba	95
England	90
Jamaica	81
South	80
China	75
Italy	73
Dominican-Republic	70
Vietnam	67
Guatemala	64

Japan	62
Poland	60
Columbia	59
Taiwan	51
Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Cambodia	19
Trinidad&Tobago	19
Laos	18
Thailand	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Honduras	13
Hungary	13
Scotland	12
Holand-Netherlands	1

Name: native_country, dtype: int64

4.8.2 Percentage of frequency distribution of values

```
[18]: for var in categorical:
      print(df[var].value_counts(normalize=True), '\n')
```

Private	0.697030
Self-emp-not-inc	0.078038
Local-gov	0.064279
?	0.056386
State-gov	0.039864
Self-emp-inc	0.034274
Federal-gov	0.029483
Without-pay	0.000430
Never-worked	0.000215

Name: workclass, dtype: float64

HS-grad	0.322502
Some-college	0.223918
Bachelors	0.164461
Masters	0.052916
Assoc-voc	0.042443

11th	0.036086
Assoc-acdm	0.032769
10th	0.028654
7th-8th	0.019840
Prof-school	0.017690
9th	0.015786
12th	0.013298
Doctorate	0.012684
5th-6th	0.010227
1st-4th	0.005160
Preschool	0.001566

Name: education, dtype: float64

Married-civ-spouse	0.459937
Never-married	0.328092
Divorced	0.136452
Separated	0.031479
Widowed	0.030497
Married-spouse-absent	0.012837
Married-AF-spouse	0.000706

Name: marital_status, dtype: float64

Prof-specialty	0.127146
Craft-repair	0.125887
Exec-managerial	0.124873
Adm-clerical	0.115783
Sales	0.112097
Other-service	0.101195
Machine-op-inspct	0.061485
?	0.056601
Transport-moving	0.049046
Handlers-cleaners	0.042075
Farming-fishing	0.030527
Tech-support	0.028500
Protective-serv	0.019932
Priv-house-serv	0.004576
Armed-Forces	0.000276

Name: occupation, dtype: float64

Husband	0.405178
Not-in-family	0.255060
Own-child	0.155646
Unmarried	0.105832
Wife	0.048156
Other-relative	0.030128

Name: relationship, dtype: float64

White	0.854274
-------	----------

Black	0.095943
Asian-Pac-Islander	0.031909
Amer-Indian-Eskimo	0.009551
Other	0.008323

Name: race, dtype: float64

Male	0.669205
Female	0.330795

Name: sex, dtype: float64

United-States	0.895857
Mexico	0.019748
?	0.017905
Philippines	0.006081
Germany	0.004207
Canada	0.003716
Puerto-Rico	0.003501
El-Salvador	0.003255
India	0.003071
Cuba	0.002918
England	0.002764
Jamaica	0.002488
South	0.002457
China	0.002303
Italy	0.002242
Dominican-Republic	0.002150
Vietnam	0.002058
Guatemala	0.001966
Japan	0.001904
Poland	0.001843
Columbia	0.001812
Taiwan	0.001566
Haiti	0.001351
Iran	0.001321
Portugal	0.001136
Nicaragua	0.001044
Peru	0.000952
France	0.000891
Greece	0.000891
Ecuador	0.000860
Ireland	0.000737
Hong	0.000614
Cambodia	0.000584
Trinidad&Tobago	0.000584
Laos	0.000553
Thailand	0.000553
Yugoslavia	0.000491
Outlying-US(Guam-USVI-etc)	0.000430

```
Honduras          0.000399
Hungary           0.000399
Scotland          0.000369
Holand-Netherlands 0.000031
Name: native_country, dtype: float64
```

Findings

- Now, we can see that there are several variables like `workclass`, `occupation` and `native_country` which contain missing values.
- Generally, the missing values are coded as `NaN` and python will detect them with the usual command of `df.isnull().sum()`.
- But, in this case the missing values are coded as `?`. Pandas fails to detect these as missing values because it does not consider `?` as missing values.
- So, we have to replace `?` with `NaN` so that Python can detect these missing values.
- We will explore these variables and replace `?` with `NaN`.

4.8.3 Explore target variable

```
[19]: # check for missing values
df['income'].isnull().sum()
```

```
[19]: 0
```

We can see that there are no missing values in the `income` target variable.

```
[20]: # view number of unique values
df['income'].nunique()
```

```
[20]: 2
```

There are 2 unique values in the `income` variable.

```
[21]: # view the unique values
df['income'].unique()
```

```
[21]: array([' <=50K', ' >50K'], dtype=object)
```

The two unique values are `<=50K` and `>50K`.

```
[22]: # view the frequency distribution of values
df['income'].value_counts()
```

```
[22]: <=50K    24720
      >50K     7841
Name: income, dtype: int64
```

```
[23]: # view percentage of frequency distribution of values
df['income'].value_counts(normalize=True)
```

```
[23]: <=50K    0.75919
      >50K    0.24081
      Name: income, dtype: float64
```


```
[24]: # visualize frequency distribution of income variable

f,ax=plt.subplots(1,2,figsize=(18,8))

ax[0] = df['income'].value_counts().plot.pie(explode=[0,0],autopct='%1.
    ↪1f%%',ax=ax[0])
ax[0].set_title('Income Share')

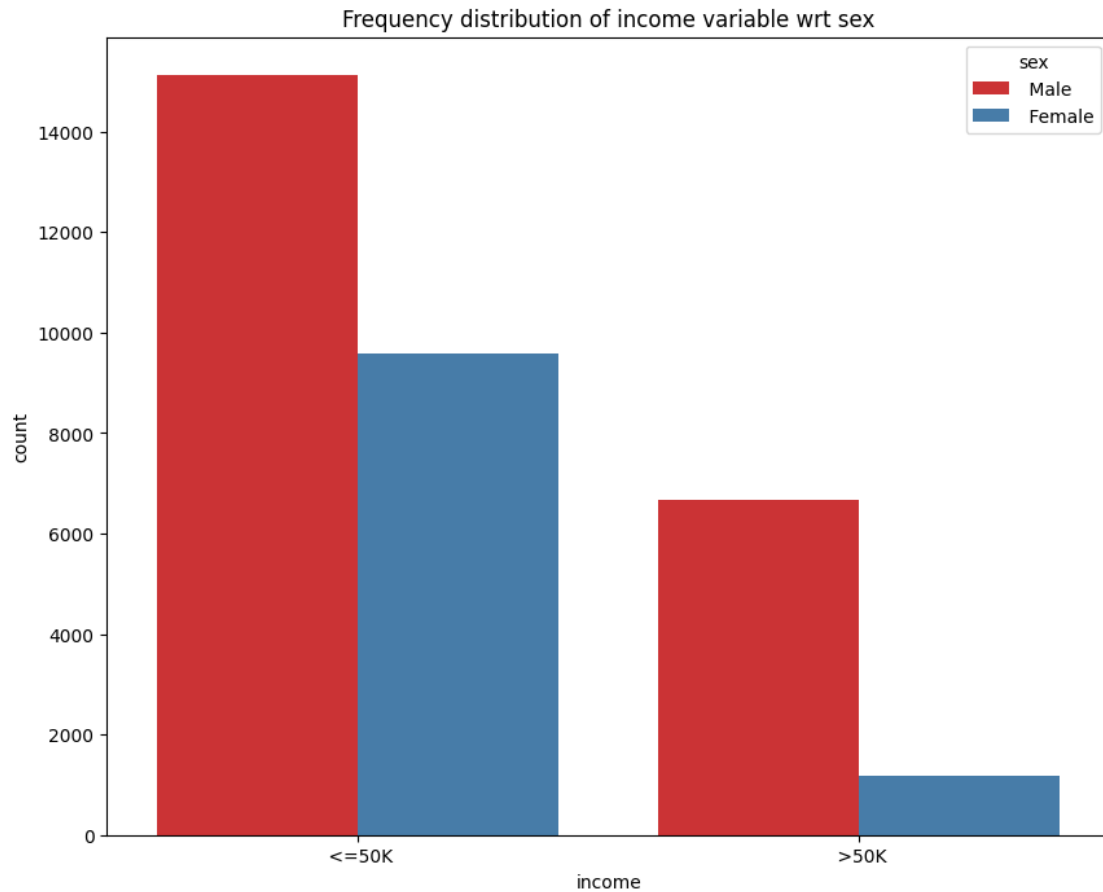
#f, ax = plt.subplots(figsize=(6, 8))
ax[1] = sns.countplot(x="income", data=df, palette="Set1")
ax[1].set_title("Frequency distribution of income variable")

plt.show()
```



4.8.4 Visualize income wrt sex variable

```
[25]: f, ax = plt.subplots(figsize=(10, 8))
      ax = sns.countplot(x="income", hue="sex", data=df, palette="Set1")
      ax.set_title("Frequency distribution of income variable wrt sex")
      plt.show()
```




Interpretation

- We can see that males make more money than females in both the income categories.

4.8.5 Visualize income wrt race

```
[26]: f, ax = plt.subplots(figsize=(10, 8))
      ax = sns.countplot(x="income", hue="race", data=df, palette="Set2")
      ax.set_title("Frequency distribution of income variable wrt race")
      plt.show()
```

Interpretation

- We can see that whites make more money than non-whites in both the income categories.

4.8.6 Explore workclass variable

```
[27]: # check number of unique labels
df.workclass.nunique()
```

[27]: 9

```
[28]: # view the unique labels
df.workclass.unique()
```

```
[28]: array([' State-gov', ' Self-emp-not-inc', ' Private', ' Federal-gov',
        ' Local-gov', ' ?', ' Self-emp-inc', ' Without-pay',
        ' Never-worked'], dtype=object)
```

```
[29]: # view frequency distribution of values
df.workclass.value_counts()
```

```
[29]: Private          22696
      Self-emp-not-inc  2541
      Local-gov        2093
      ?                1836
      State-gov        1298
      Self-emp-inc      1116
      Federal-gov       960
      Without-pay       14
      Never-worked       7
      Name: workclass, dtype: int64
```

We can see that there are 1836 values encoded as ? in workclass variable. I will replace these ? with NaN.

```
[30]: # replace '?' values in workclass variable with `NaN`
      df['workclass'].replace('?', np.NaN, inplace=True)
```

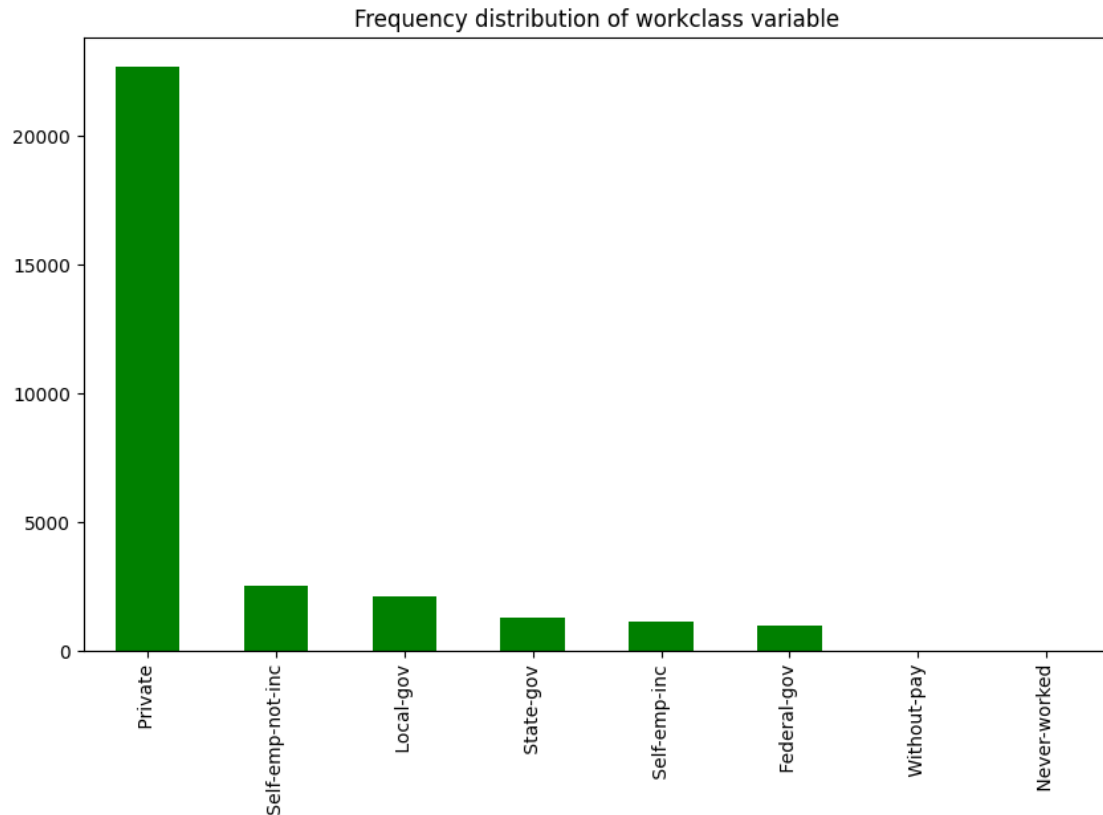
```
[31]: # again check the frequency distribution of values in workclass variable
      df.workclass.value_counts()
```

```
[31]: Private          22696
      Self-emp-not-inc  2541
      Local-gov        2093
      State-gov        1298
      Self-emp-inc      1116
      Federal-gov       960
      Without-pay       14
      Never-worked       7
      Name: workclass, dtype: int64
```

- Now, we can see that there are no values encoded as ? in the workclass variable.
- We will adopt similar approach with occupation and native_country column.

4.8.7 Visualize workclass variable

```
[32]: f, ax = plt.subplots(figsize=(10, 6))
      ax = df.workclass.value_counts().plot(kind="bar", color="green")
      ax.set_title("Frequency distribution of workclass variable")
      ax.set_xticklabels(df.workclass.value_counts().index, rotation=90)
      plt.show()
```




Interpretation

- We can see that there are lot more private workers than other category of workers.

4.8.8 Visualize workclass variable wrt income variable

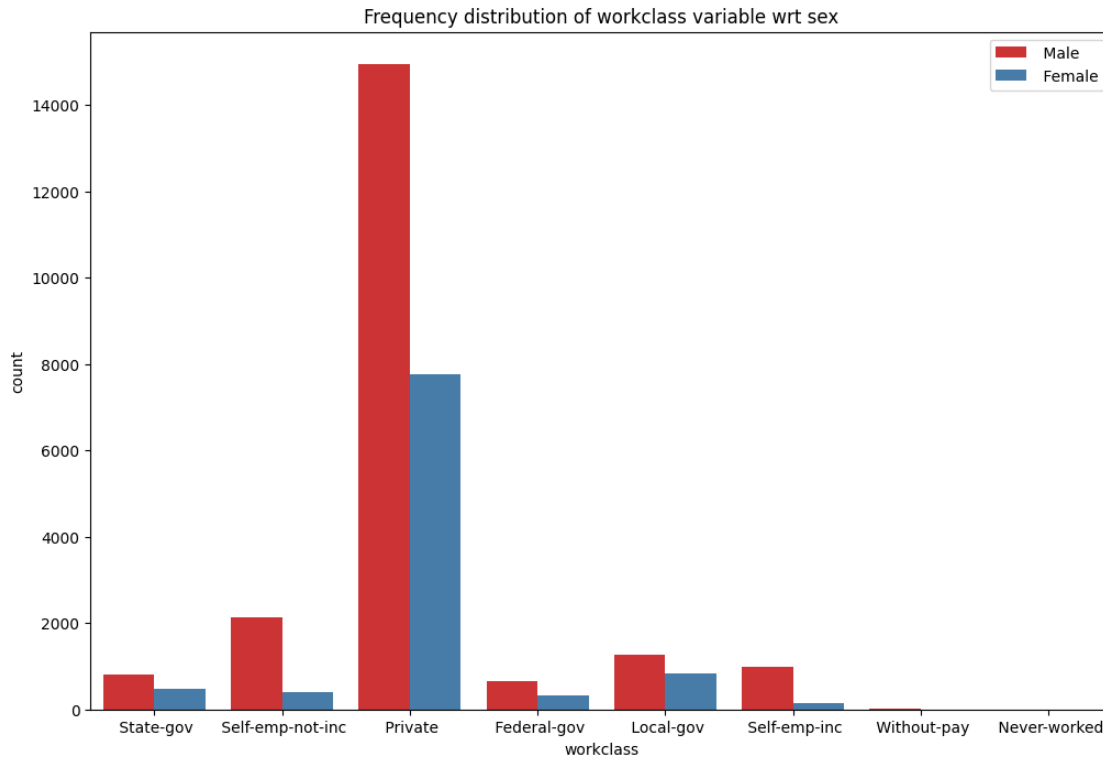
```
[33]: f, ax = plt.subplots(figsize=(12, 8))
      ax = sns.countplot(x="workclass", hue="income", data=df, palette="Set1")
      ax.set_title("Frequency distribution of workclass variable wrt income")
      ax.legend(loc='upper right')
      plt.show()
```



Interpretation - We can see that workers make less than equal to 50k in most of the working categories. - But this trend is more appealing in Private `workclass` category.

4.8.9 Visualize workclass variable wrt sex variable

```
[34]: f, ax = plt.subplots(figsize=(12, 8))
      ax = sns.countplot(x="workclass", hue="sex", data=df, palette="Set1")
      ax.set_title("Frequency distribution of workclass variable wrt sex")
      ax.legend(loc='upper right')
      plt.show()
```



Interpretation - We can see that there are more male workers than female workers in all the working category. - The trend is more appealing in Private sector.

4.8.10 Explore occupation variable

```
[35]: # check number of unique labels
df.occupation.nunique()
```

[35]: 15

```
[36]: # view unique labels
df.occupation.unique()
```

```
[36]: array([' Adm-clerical', ' Exec-managerial', ' Handlers-cleaners',
        ' Prof-specialty', ' Other-service', ' Sales', ' Craft-repair',
        ' Transport-moving', ' Farming-fishing', ' Machine-op-inspct',
        ' Tech-support', ' ?', ' Protective-serv', ' Armed-Forces',
        ' Priv-house-serv'], dtype=object)
```

```
[37]: # view frequency distribution of values
df.occupation.value_counts()
```

```
[37]: Prof-specialty      4140
      Craft-repair      4099
      Exec-managerial   4066
      Adm-clerical      3770
      Sales             3650
      Other-service     3295
      Machine-op-inspct 2002
      ?                1843
      Transport-moving  1597
      Handlers-cleaners 1370
      Farming-fishing   994
      Tech-support      928
      Protective-serv   649
      Priv-house-serv   149
      Armed-Forces      9
      Name: occupation, dtype: int64
```


We can see that there are 1843 values encoded as ? in occupation variable. I will replace these ? with NaN.

```
[38]: # replace '?' values in occupation variable with `NaN`
      df['occupation'].replace('?', np.NaN, inplace=True)
```

```
[39]: # again check the frequency distribution of values
      df.occupation.value_counts()
```

```
[39]: Prof-specialty      4140
      Craft-repair      4099
      Exec-managerial   4066
      Adm-clerical      3770
      Sales             3650
      Other-service     3295
      Machine-op-inspct 2002
      Transport-moving  1597
      Handlers-cleaners 1370
      Farming-fishing   994
      Tech-support      928
      Protective-serv   649
      Priv-house-serv   149
      Armed-Forces      9
      Name: occupation, dtype: int64
```

```
[40]: # visualize frequency distribution of `occupation` variable
      f, ax = plt.subplots(figsize=(12, 8))
      ax = sns.countplot(x="occupation", data=df, palette="Set1")
      ax.set_title("Frequency distribution of occupation variable")
      ax.set_xticklabels(df.occupation.value_counts().index, rotation=90)
      plt.show()
```



4.8.11 Explore native_country variable

```
[41]: # check number of unique labels
df.native_country.nunique()
```

[41]: 42

```
[42]: # view unique labels
df.native_country.unique()
```

```
[42]: array([' United-States', ' Cuba', ' Jamaica', ' India', ' ?', ' Mexico',
        ' South', ' Puerto-Rico', ' Honduras', ' England', ' Canada',
        ' Germany', ' Iran', ' Philippines', ' Italy', ' Poland',
        ' Columbia', ' Cambodia', ' Thailand', ' Ecuador', ' Laos',
        ' Taiwan', ' Haiti', ' Portugal', ' Dominican-Republic',
        ' El-Salvador', ' France', ' Guatemala', ' China', ' Japan',
        ' Yugoslavia', ' Peru', ' Outlying-US(Guam-USVI-etc)', ' Scotland',
        ' Trinidad&Tobago', ' Greece', ' Nicaragua', ' Vietnam', ' Hong',
```

```
' Ireland', ' Hungary', ' Holand-Netherlands'], dtype=object)
```

```
[43]: # check frequency distribution of values
df.native_country.value_counts()
```

```
[43]: United-States      29170
      Mexico          643
      ?              583
      Philippines     198
      Germany         137
      Canada          121
      Puerto-Rico     114
      El-Salvador     106
      India           100
      Cuba            95
      England         90
      Jamaica         81
      South           80
      China           75
      Italy           73
      Dominican-Republic 70
      Vietnam         67
      Guatemala       64
      Japan           62
      Poland          60
      Columbia        59
      Taiwan          51
      Haiti           44
      Iran            43
      Portugal        37
      Nicaragua       34
      Peru            31
      France          29
      Greece          29
      Ecuador         28
      Ireland         24
      Hong            20
      Cambodia        19
      Trinidad&Tobago 19
      Laos            18
      Thailand        18
      Yugoslavia      16
      Outlying-US(Guam-USVI-etc) 14
      Honduras        13
      Hungary         13
      Scotland        12
      Holand-Netherlands 1
```


Name: native_country, dtype: int64

We can see that there are 583 values encoded as ? in native_country variable. I will replace these ? with NaN.

```
[44]: # replace '?' values in native_country variable with `NaN`  
df['native_country'].replace('?', np.NaN, inplace=True)
```


```
[45]: # again check the frequency distribution of values  
df.native_country.value_counts()
```

```
[45]: United-States      29170  
      Mexico           643  
      Philippines      198  
      Germany          137  
      Canada           121  
      Puerto-Rico      114  
      El-Salvador      106  
      India            100  
      Cuba             95  
      England          90  
      Jamaica          81  
      South            80  
      China            75  
      Italy            73  
      Dominican-Republic 70  
      Vietnam          67  
      Guatemala        64  
      Japan            62  
      Poland           60  
      Columbia         59  
      Taiwan           51  
      Haiti            44  
      Iran             43  
      Portugal         37  
      Nicaragua        34  
      Peru             31  
      France           29  
      Greece           29  
      Ecuador          28  
      Ireland          24  
      Hong             20  
      Cambodia         19  
      Trinidad&Tobago  19  
      Laos             18  
      Thailand         18  
      Yugoslavia       16
```

Outlying-US(Guam-USVI-etc)	14
Honduras	13
Hungary	13
Scotland	12
Holand-Netherlands	1

Name: native_country, dtype: int64

```
[46]: # visualize frequency distribution of `native_country` variable
f, ax = plt.subplots(figsize=(16, 12))
ax = sns.countplot(x="native_country", data=df, palette="Set1")
ax.set_title("Frequency distribution of native_country variable")
ax.set_xticklabels(df.native_country.value_counts().index, rotation=90)
plt.show()
```



We can see that United-States dominate amongst the `native_country` variables.

4.8.12 Check missing values in categorical variables

```
[47]: df[categorical].isnull().sum()
```

```
[47]: workclass      1836
      education      0
      marital_status  0
      occupation     1843
      relationship    0
      race            0
      sex            0
      native_country   583
      dtype: int64
```

Now, we can see that `workclass`, `occupation` and `native_country` variable contains missing values.

4.8.13 Number of labels: Cardinality

- The number of labels within a categorical variable is known as **cardinality**.
- A high number of labels within a variable is known as **high cardinality**.
- High cardinality may pose some serious problems in the machine learning model. So, we will check for high cardinality.

```
[48]: # check for cardinality in categorical variables
      for var in categorical:
          print(var, ' contains ', len(df[var].unique()), ' labels')
```

```
workclass contains 9 labels
education contains 16 labels
marital_status contains 7 labels
occupation contains 15 labels
relationship contains 6 labels
race contains 5 labels
sex contains 2 labels
native_country contains 42 labels
```

We can see that `native_country` column contains relatively large number of labels as compared to other columns.

4.9 Explore Numerical Variables

```
[49]: numerical = ['age', 'fnlwgt', 'education_num', 'capital_gain',
                  'capital_loss', 'hours_per_week']
```

4.9.1 Preview the numerical variables

```
[50]: df[numerical].head()
```

```
[50]:   age  fnlwgt  education_num  capital_gain  capital_loss  hours_per_week
0   39   77516             13         2174           0             40
1   50   83311             13           0           0             13
2   38  215646              9           0           0             40
3   53  234721              7           0           0             40
4   28  338409             13           0           0             40
```

4.9.2 Check missing values in numerical variables


```
[51]: df[numerical].isnull().sum()
```

```
[51]: age                0
      fnlwgt            0
      education_num     0
      capital_gain      0
      capital_loss      0
      hours_per_week    0
      dtype: int64
```

We can see that there are no missing values in the numerical variables.

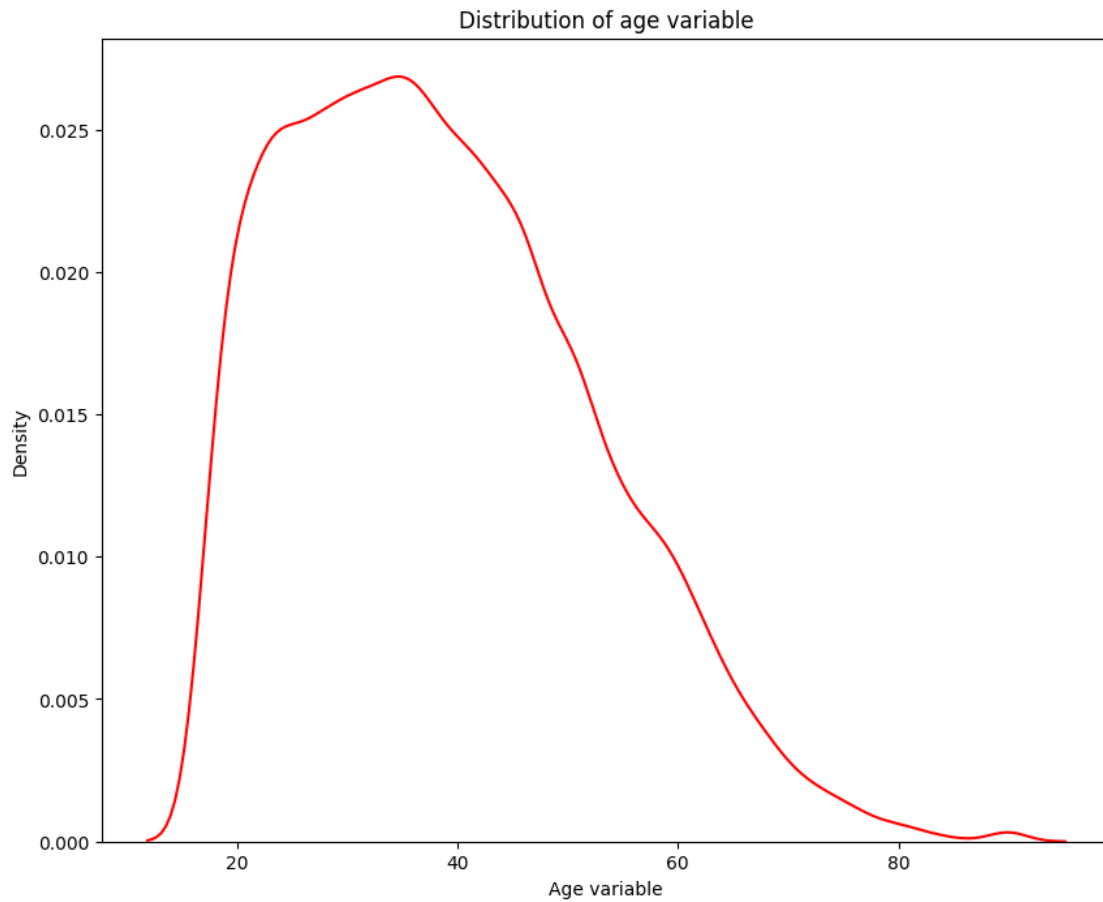
4.9.3 Explore age variable

```
[52]: # View the distribution of `age` variable
      f, ax = plt.subplots(figsize=(10,8))
      x = df['age']
      ax = sns.histplot(x, bins=10, color='blue', kde=True)
      ax.set_title("Distribution of age variable")
      plt.show()
```



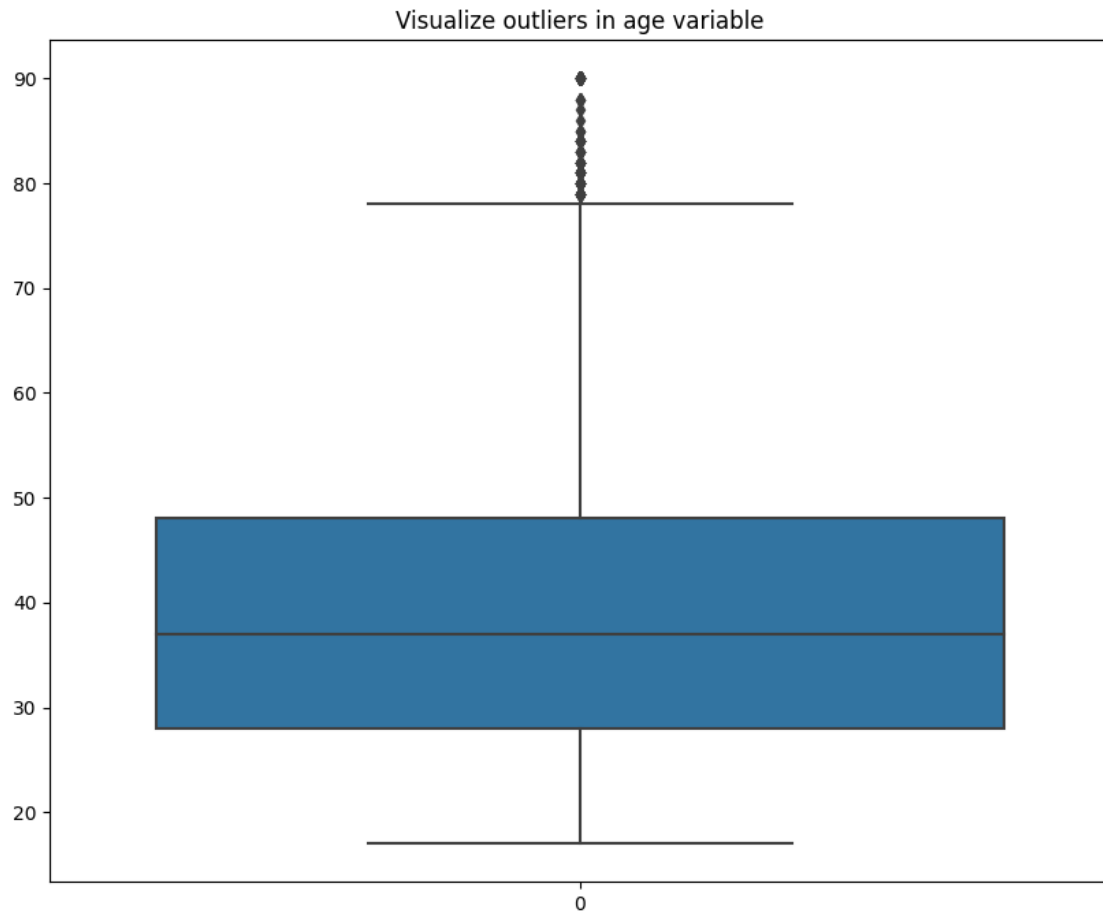
We can see that `age` is slightly positively skewed.

```
[53]: f, ax = plt.subplots(figsize=(10,8))
x = df['age']
x = pd.Series(x, name="Age variable")
ax = sns.kdeplot(x, color='red')
ax.set_title("Distribution of age variable")
plt.show()
```



4.9.4 Detect outliers in age variable with boxplot

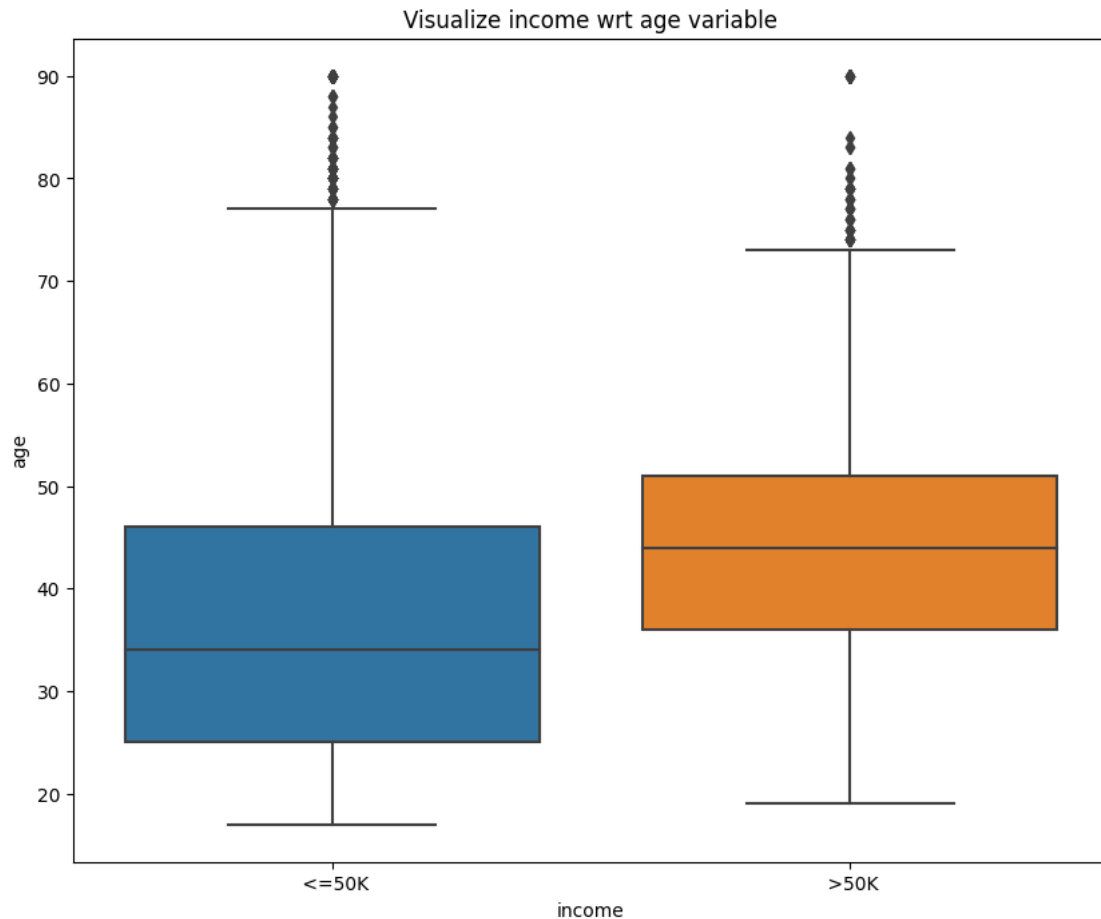
```
[54]: f, ax = plt.subplots(figsize=(10,8))
      x = df['age']
      ax = sns.boxplot(x)
      ax.set_title("Visualize outliers in age variable")
      plt.show()
```



We can see that there are lots of outliers in **age** variable.

4.9.5 Explore relationship between age and income variables

```
[55]: f, ax = plt.subplots(figsize=(10, 8))
      ax = sns.boxplot(x="income", y="age", data=df)
      ax.set_title("Visualize income wrt age variable")
      plt.show()
```




Interpretation

- As expected, younger people make less money as compared to senior people.


4.9.6 Visualize income wrt age and sex variable

```
[56]: f, ax = plt.subplots(figsize=(10, 8))
      ax = sns.boxplot(x="income", y="age", hue="sex", data=df)
      ax.set_title("Visualize income wrt age and sex variable")
      ax.legend(loc='upper right')
      plt.show()
```

```
[57]: plt.figure(figsize=(8,6))
      ax = sns.catplot(x="income", y="age", col="sex", data=df, kind="box", height=8,
      ↳ aspect=1)
      plt.show()
```

<Figure size 800x600 with 0 Axes>

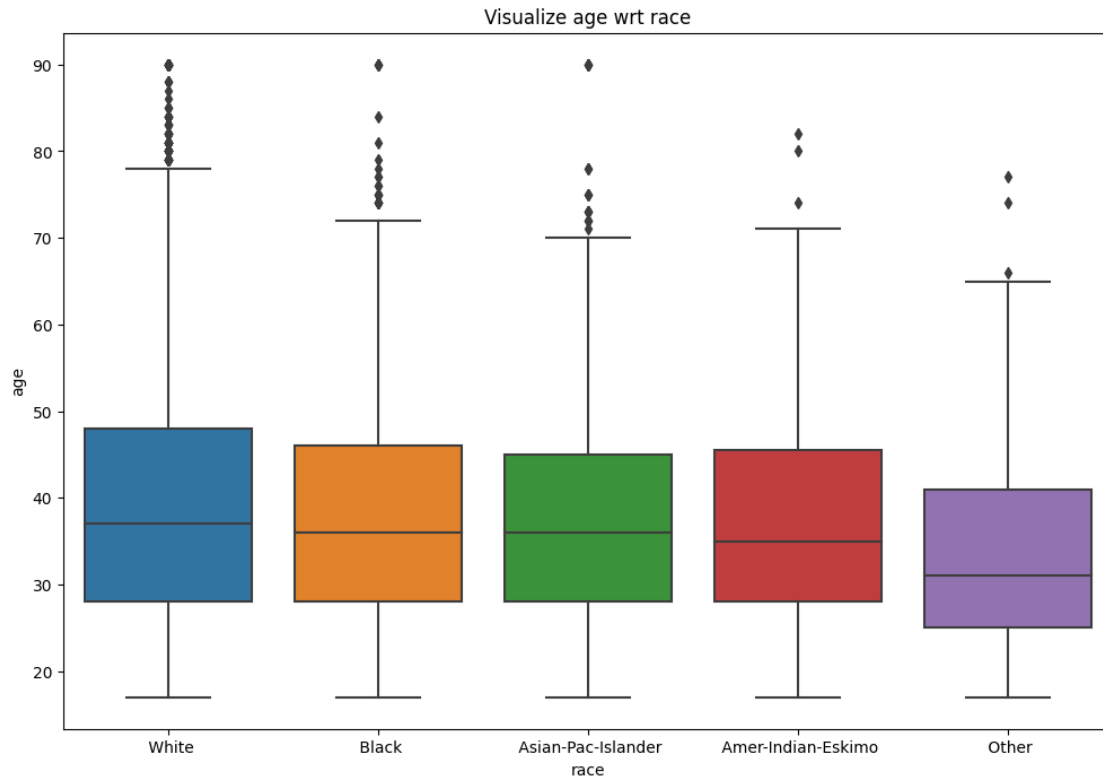


Interpretation

- Senior people make more money than younger people.

4.9.7 Visualize relationship between race and age

```
[58]: plt.figure(figsize=(12,8))
sns.boxplot(x='race', y="age", data = df)
plt.title("Visualize age wrt race")
plt.show()
```



Interpretation

- Whites are more older than other groups of people.

4.9.8 Find out the correlations

```
[59]: df.corr() # Compute pairwise correlation of columns, excluding NA/null values.
```

/tmp/ipykernel_107355/605702514.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
df.corr() # Compute pairwise correlation of columns, excluding NA/null values.
```

```
[59]:
```

	age	fnlwgt	education_num	capital_gain	capital_loss	\
age	1.000000	-0.076646	0.036527	0.077674	0.057775	
fnlwgt	-0.076646	1.000000	-0.043195	0.000432	-0.010252	
education_num	0.036527	-0.043195	1.000000	0.122630	0.079923	
capital_gain	0.077674	0.000432	0.122630	1.000000	-0.031615	
capital_loss	0.057775	-0.010252	0.079923	-0.031615	1.000000	
hours_per_week	0.068756	-0.018768	0.148123	0.078409	0.054256	

hours_per_week

```

age                0.068756
fnlwgt             -0.018768
education_num      0.148123
capital_gain       0.078409
capital_loss       0.054256
hours_per_week     1.000000

```

Interpretation - We can see that there is no strong correlation between variables.

```

[60]: # plot correlation heatmap to find out correlations
fig = plt.figure(dpi=300)
sns.heatmap(df.corr(),annot=True,linewidths=.5,cmap="viridis")
plt.show()

```

/tmp/ipykernel_107355/1293130057.py:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sns.heatmap(df.corr(),annot=True,linewidths=.5,cmap="viridis")
```

